

TP2: TeaStore Performance Efficiency

Group 02

Chun-An Bau
2165883
chun-an.bau@polymtl.ca

Youssef Amine BenDiha
2113648
youssef-amine.ben-diha@polymtl.ca

Jakub Profota
2165556
jakub.profota@polymtl.ca

CONTENTS

I	Introduction	2
II	Quality Plan	2
II-A	Quality Goals	2
II-B	Quality Assurance Strategies	2
III	Deployment	2
III-A	Method 1	3
III-B	Method 2	3
IV	Performance Profiling	3
IV-A	Overall Hot Spots	3
IV-B	Test Case Hot Spots	5
V	Version 1.0.1	5
V-A	20 Users	6
V-B	200 Users	7
V-C	500 Users	8
V-D	1000 Users	9
VI	Current Version	10
VI-A	20 Users	10
VI-B	200 Users	11
VI-C	500 Users	12
VI-D	1000 Users	13
VII	Discussion	14
VII-A	Time behavior	14
VII-B	Resource utilization	14
VII-C	Capacity	14
	References	14

Abstract

This document further elaborates on the quality plan for one of the subsystems of TeaStore [1], focusing on the Performance quality characteristic of ISO 25010 standard. The latest version of TeaStore is evaluated based on presented quality metrics and compared with previous versions.

I. INTRODUCTION

In this paper, we continue evaluating our quality assurance plan on the Authentication service of the TeaStore, this time focusing on one specific quality characteristic of ISO 25010: Performance. The Authentication service performance efficiency is crucial as this service is responsible for user account logins and logouts. These accounts contain purchase history. Not being able to keep up with an increased load may cause severe issues like ignoring user purchases. Therefore, extensive system testing under load is mandatory.

II. QUALITY PLAN

A. Quality Goals

The specified quality characteristics and their respective sub-characteristics of the Authentication subsystem are presented in Table I. We updated the table with the newly evaluated Performance characteristic.

In the Performance quality characteristic, we focused on Time behavior, Resource utilization, and Capacity sub-characteristics. Time behavior is measured by response times of incoming user requests. We believe the maximum waiting time should be no more than one second. Resource utilization is evaluated through monitoring CPU, RAM, and disk usage. As all the testing was conducted on a minimal virtual machine with reduced resources, we set the objective to not use any of the three resources at 100% capacity throughout the whole benchmark. Finally, for the Capacity sub-characteristic, we measure the ratio of successful and unsuccessful user requests. As the Authentication service is crucial, we believe at least 99% of user requests must pass.

TABLE I
QUALITY GOALS OF THE AUTHENTICATION SUBSYSTEM

Quality characteristics	Sub-characteristics	Quality measure	Objective
Functional Suitability	Completeness	Compliance with the specifications	All the specified functions of the system must be present
	Correctness	Unit Tests	All tests must pass
Reliability	Maturity	Delay mean time	Meantime 3 seconds maximum
	Availability	Downtime	Downtime 1% maximum under load
Maintainability	Modularity	Modularity index	Index reaches 0.5
	Testability	Code coverage	Test cases reach 70% of code coverage
Performance	Time behavior	Response time	Average less than 2.5 seconds
	Resource utilization	Available resources	Utilization less than 100% all the time
	Capacity	Number of successful requests	At least 99% requests passed

B. Quality Assurance Strategies

The Performance characteristic is always measured on the final product. Even though there may be testings throughout the development process, what matters is the final product deployment. Therefore, all the newly introduced sub-characteristics in the Table II should be evaluated at the system testing stage. Objective and consistent testing is usually conducted by third parties as any internal performance testing may be biased.

TABLE II
QUALITY ASSURANCE STRATEGIES OF THE AUTHENTICATION SUBSYSTEM

Scope	Stage	Roles involved
Completeness	System testing	Clients
Correctness	Development	Developers
Maturity	System testing	Third party
Availability	System testing	Third party
Modularity	Development	Developers
Testability	Development	Developers
Time behavior	System testing	Third party
Resource utilization	System testing	Third party
Capacity	System testing	Third party

III. DEPLOYMENT

There are two methods on how to deploy a Docker image of an older version of TeaStore:

A. Method 1

To deploy the reduced version of the Docker container, we have to change the version number indicated by lines 5 7 of POM.XML. The latest version is 1.4.0, and we changed it to 1.0.1 for the following experiments.

Fig. 1. pom.xml

```
5 <properties>
6 |   <teastoreversion>1.4.0</teastoreversion>
7 </properties>
```

After modification, we can launch the reduced version Docker image by the commands below.

```
$ docker run -p 3306:3306 -d --name teastore-db descartesresearch/teastore-db
$ docker run -e "DB_HOST=teastore-db" -p 8080:8080 -d --link
  teastore-db:teastore-db --name teastore-all descartesresearch/teastore-all
```

B. Method 2

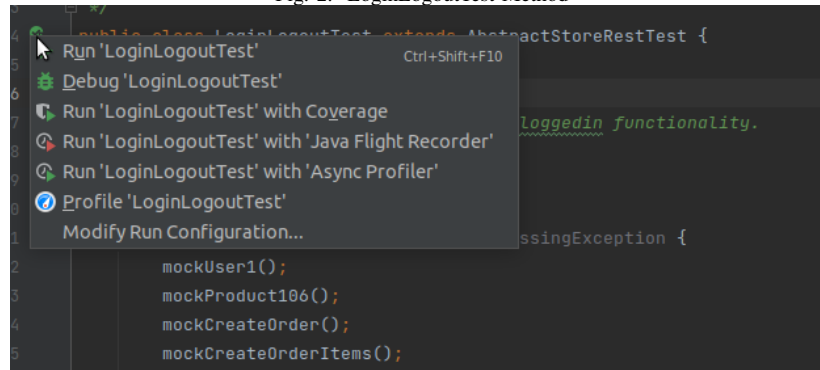
To deploy the reduced version of the docker container, we specify the version while executing the commands. For example, we launch version 1.0.1 by the commands below.

```
$ docker run -p 3306:3306 -d --name teastore-db
  descartesresearch/teastore-db:1.0.1
$ docker run -e "DB_HOST=teastore-db" -p 8080:8080 -d --link
  teastore-db:teastore-db --name teastore-all
  descartesresearch/teastore-all:1.0.1
```

IV. PERFORMANCE PROFILING

In version v1.0.1, we found a test case LOGINLOGOUTTEST.JAVA, which contains the tests of logging in/out, adding items to the cart, removing items from the cart, and placing the order. We profiled by IntelliJ's plugin, JProfiler [2].

Fig. 2. LoginLogoutTest Method



After selecting *Profile LoginLogoutTest*, the JProfiler UI session will pop up, and we choose the calling methods *instrumentation* and *full sampling* with default settings for experiments as seen on Fig. 3.

A. Overall Hot Spots

The aggregation level of the hot spot results is class level. The first diagram is instrumentation, and the second is full sampling. We found that class `JAVA.UTIL.CONCURRENT.THREADPOOLEXECUTOR$WRAPPINGRUNNABLE` occupies the most time slots in both methods, while some classes have significant measure time between the two methods. For example, `ORG.APACHE.TOMCAT.UTIL.THREADS.TASKTHREAD$WRAPPINGRUNNABLE` occupied 23% of measure times in instrumentation but only occupied less than 1% in full sampling. Also, the result shows that instrumentation takes more time than full sampling, which matches our expectations. Results of instrumentation and full sampling are depicted on Fig. 4 and Fig. 5, respectively.

Fig. 3. Session Startup

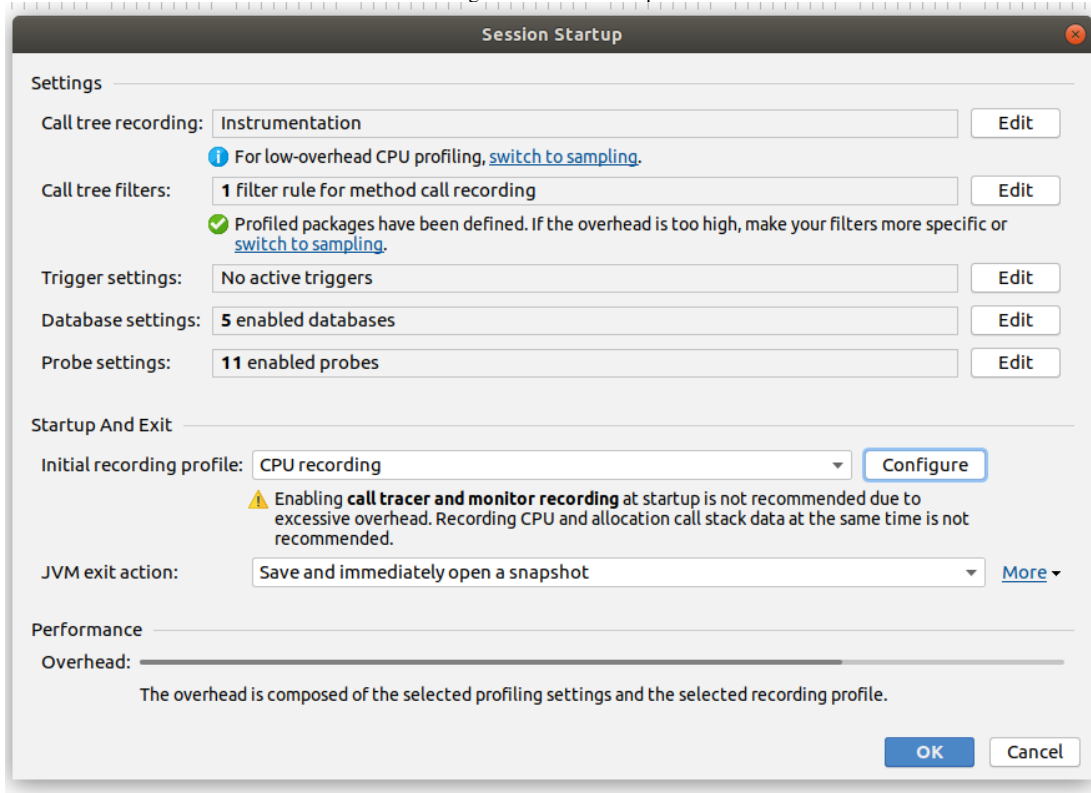
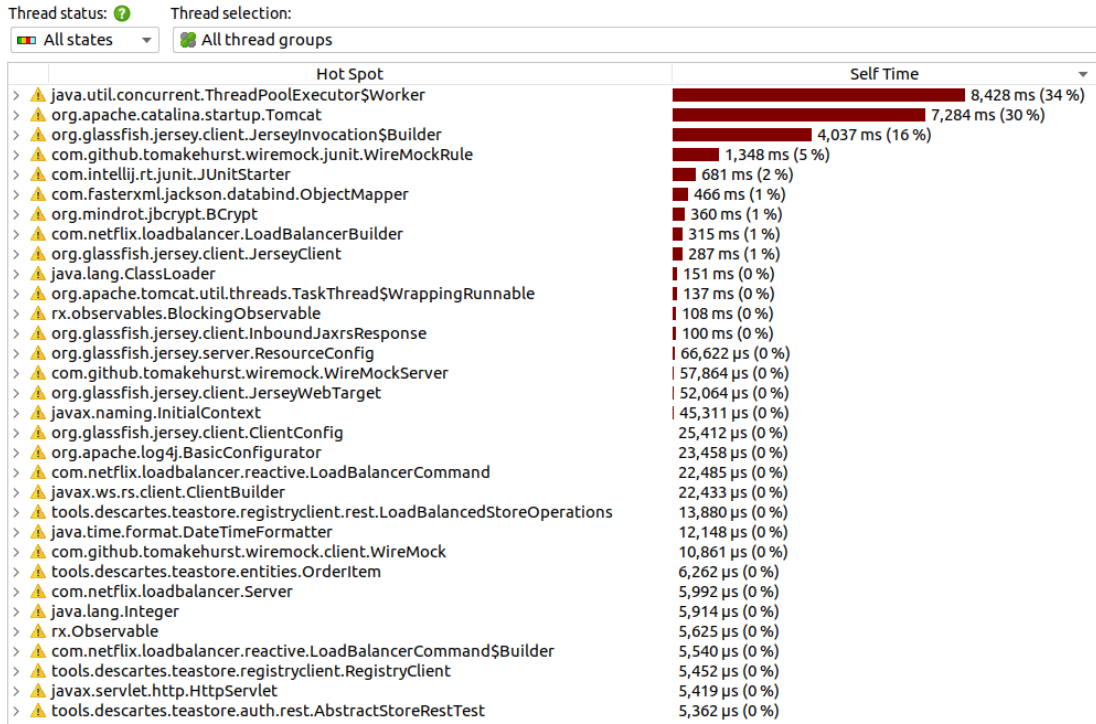


Fig. 4. Instrumentation Results

Thread status: ?		Thread selection:	
<input type="text" value="All states"/>		<input type="text" value="All thread groups"/>	
		Hot Spot	
		Self Time	
>	java.util.concurrent.ThreadPoolExecutor\$Worker	13,212 ms (36 %)	
>	org.apache.tomcat.util.threads.TaskThread\$WrappingRunnable	8,365 ms (23 %)	
>	org.apache.catalina.startup.Tomcat	6,971 ms (19 %)	
>	javax.ws.rs.client.Invocation\$Builder	3,616 ms (10 %)	
>	com.github.tomakehurst.wiremock.junit.WireMockRule	1,365 ms (3 %)	
>	com.intellij.rt.junit.JUnit4TestRunner	352 ms (0 %)	
>	org.mindrot.jbcrypt.BCrypt	322 ms (0 %)	
>	javax.ws.rs.client.Client	305 ms (0 %)	
>	tools.descartes.teastore.auth.rest.AbstractStoreRestTest	294 ms (0 %)	
>	com.fasterxml.jackson.databind.ObjectMapper	291 ms (0 %)	
>	com.netflix.loadbalancer.LoadBalancerBuilder	244 ms (0 %)	
>	javax.ws.rs.core.Response	104 ms (0 %)	
>	rx.observables.BlockingObservable	80,247 μs (0 %)	
>	org.glassfish.jersey.server.ResourceConfig	62,337 μs (0 %)	
>	javax.naming.InitialContext	44,934 μs (0 %)	
>	javax.ws.rs.client.WebTarget	44,067 μs (0 %)	
>	org.apache.log4j.BasicConfigurator	28,553 μs (0 %)	
>	com.netflix.loadbalancer.reactive.LoadBalancerCommand	28,493 μs (0 %)	
>	javax.ws.rs.client.ClientBuilder	27,060 μs (0 %)	
>	org.glassfish.jersey.client.ClientConfig	21,829 μs (0 %)	
>	tools.descartes.teastore.registryclient.loadbalancers.ServiceLoadBalancer	17,899 μs (0 %)	
>	tools.descartes.teastore.auth.rest.AuthUserActionsREST	17,867 μs (0 %)	
>	tools.descartes.teastore.registryclient.rest.LoadBalancedStoreOperations	13,814 μs (0 %)	
>	tools.descartes.teastore.auth.rest.EmptyAuthStartup	11,536 μs (0 %)	
>	com.github.tomakehurst.wiremock.client.WireMock	7,692 μs (0 %)	
>	java.security.MessageDigest	5,522 μs (0 %)	
>	tools.descartes.teastore.registryclient.RegistryClient	4,988 μs (0 %)	
>	tools.descartes.teastore.registryclient.rest.LoadBalancedCRUDOperations	4,553 μs (0 %)	
>	tools.descartes.teastore.auth.rest.LoginLogoutTest	4,161 μs (0 %)	
>	java.net.URLEncoder	4,044 μs (0 %)	
>	rx.Observable	3,321 μs (0 %)	
>	tools.descartes.teastore.entities.Order	3,263 μs (0 %)	


Fig. 5. Full Sampling Results



B. Test Case Hot Spots

Fig. 6 shows the result while we only focus on the performance of the test case at the method level.

Fig. 6. Method Level Results



We can see that the period `ISLOGGEDIN()` takes 10-100 times more than other operations, such as `LOGIN()`, `LOGOUT()`, and `PLACEORDER()`. If we look at the code snippet, we will find that the method calls external services via REST API, which needs more time to complete.

Fig. 7. isLodedIn() method


```
public static boolean isLoggedIn(SessionBlob blob) throws NotFoundException, LoadBalancerTimeoutException {
    Response r = ServiceLoadBalancer.loadBalanceRESTOperation(Service.AUTH, endpointURL: "useractions", Product.class,
        client -> ResponseWrapper.wrap(HttpWrapper.wrap(client.getEndpointTarget().path("isLoggedIn"))
            .post(Entity.entity(blob, MediaType.APPLICATION_JSON), Response.class)));
    return RestUtil.readThrowAndOrClose(r, SessionBlob.class) != null;
}
```

If the developers want to improve the overhead, they may use an alternative way to communicate externally. However, it might sacrifice the convenience of maintaining the project.



V. VERSION 1.0.1

Load testing was conducted on a virtual machine. We used TeaStore version 1.0.1, which was deployed locally as a Docker image. The results of the current version follow in the next section. We simulated a load of 20, 200, 500, and 1000 users. The following diagrams are as follows: resource utilization throughout the benchmark, response times of methods, cumulative number of requests in four response time categories, and finally pass/fail ratio of user requests.


A. 20 Users




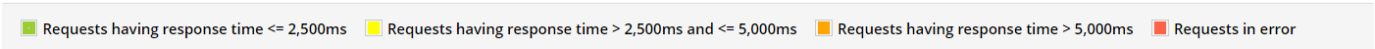

Requests	Executions			Response Times (ms)								Throughput	Network (KB/sec)	
	Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total		800	0	0.00%	76.47	5	367	63.00	147.00	173.80	244.96	214.88	1955.75	38.35
HomePage		200	0	0.00%	83.93	5	367	69.00	152.90	183.95	346.99	54.07	465.81	7.87
LoggedIn		200	0	0.00%	72.49	7	307	58.00	142.90	173.55	244.92	56.85	489.78	14.77
LoginPage		200	0	0.00%	76.07	5	264	67.00	141.80	173.80	246.86	56.34	540.16	8.47
Logout		200	0	0.00%	73.39	6	241	63.50	141.70	162.95	225.98	57.41	550.40	9.08




B. 200 Users






Requests	Executions				Response Times (ms)							Throughput	Network (KB/sec)	
	Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total		4000	712	17.80%	6761.38	30	40960	888.00	23104.10	26420.80	34607.00	29.45	236.63	5.26
HomePage		1000	74	7.40%	7618.52	271	26365	1010.00	23094.60	23219.85	24242.87	7.38	59.75	1.07
LoggedIn		1000	252	25.20%	7045.82	30	35116	850.50	23969.70	26394.95	34483.09	8.65	63.61	2.25
LoginPage		1000	196	19.60%	5869.84	104	35190	949.50	11601.80	25854.90	34744.87	7.90	65.28	1.19
Logout		1000	190	19.00%	6511.34	30	40960	726.00	33872.50	34246.05	40863.41	9.56	80.55	1.51




C. 500 Users






Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total	6000	5924	98.73%	51316.37	8761	76261	52590.00	64631.90	66766.95	70145.94	9.45	54.60	1.69
HomePage	1500	1500	100.00%	43347.29	15507	58281	44865.00	54767.10	56089.60	57319.17	2.91	16.68	0.42
LoggedIn	1500	1500	100.00%	47458.14	33397	61378	46788.00	56348.80	58043.50	59822.83	2.68	15.68	0.70
LoginPage	1500	1500	100.00%	58701.45	44754	73421	58102.50	67241.60	68507.00	70491.84	2.68	15.39	0.40
Logout	1500	1424	94.93%	55758.60	8761	76261	55604.50	66104.00	67433.25	74443.29	2.88	16.69	0.46



D. 1000 Users




Requests	Executions			Response Times (ms)								Throughput	Network (KB/sec)	
	Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total		4000	4000	100.00%	81252.19	0	208541	74496.00	180377.80	183651.10	201425.81	12.18	52.86	1.40
HomePage		1000	1000	100.00%	34141.66	10003	59953	36385.50	57412.50	58498.70	59904.97	16.37	75.71	2.38
LogedIn		1000	1000	100.00%	140896.46	82544	208541	169388.50	197006.00	201345.85	201590.99	3.92	17.61	0.64
LoginPage		1000	1000	100.00%	75760.45	61100	104210	74673.00	88426.80	90419.65	95887.39	6.45	37.08	0.97
Logout		1000	1000	100.00%	74210.19	0	171058	32425.00	161297.50	167268.55	170732.95	5.79	14.45	0.00


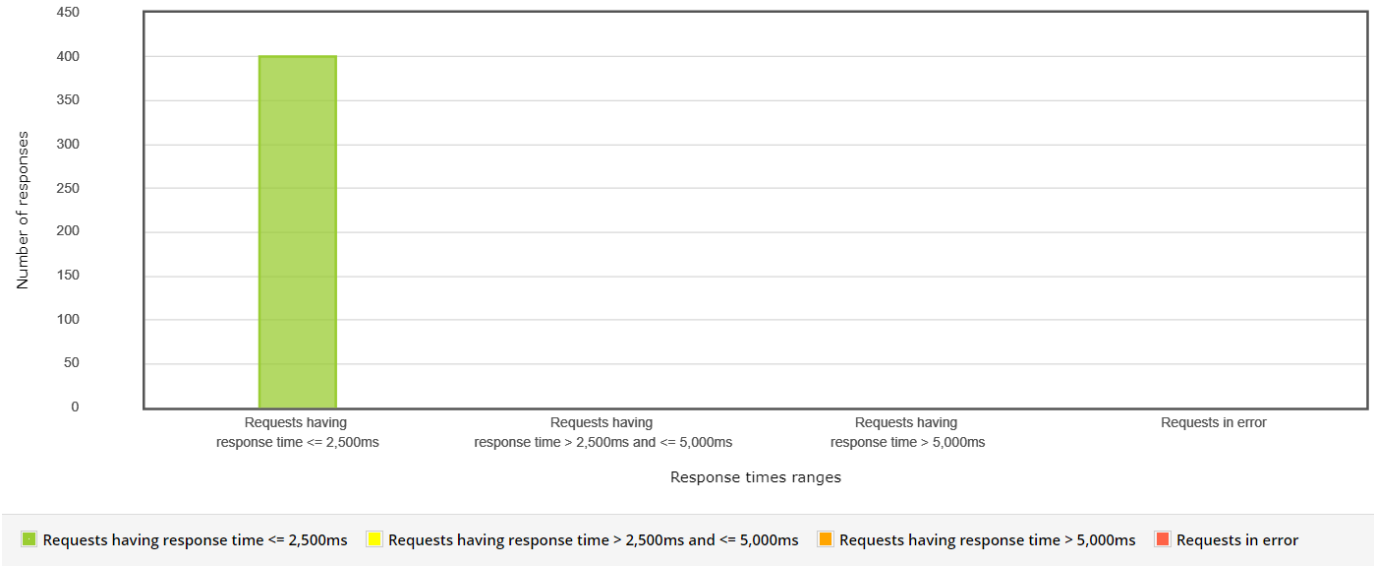


VI. CURRENT VERSION


A. 20 Users





Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total	400	0	0.00%	429.81	15	1647	311.00	828.50	1631.50	1640.00	43.56	396.45	7.77
HomePage	100	0	0.00%	660.52	30	1647	361.50	1635.00	1639.95	1646.93	11.03	95.05	1.61
LoggedIn	100	0	0.00%	362.42	23	977	305.00	748.10	787.75	976.81	14.71	126.73	3.82
LoginPage	100	0	0.00%	370.73	15	1377	322.00	609.80	721.40	1373.00	13.50	129.43	2.03
Logout	100	0	0.00%	325.58	17	1290	268.50	672.00	844.30	1289.19	15.35	147.17	2.43




B. 200 Users




Requests	Executions			Response Times (ms)								Throughput	Network (KB/sec)	
	Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total		4000	0	0.00%	4229.27	15	67315	1129.00	2680.60	48764.20	66527.90	42.85	386.37	7.65
HomePage		1000	0	0.00%	13106.84	31	67315	1378.00	65778.50	66450.80	66829.99	10.72	88.69	1.56
LoggedIn		1000	0	0.00%	1179.13	15	5042	1003.50	2187.50	2498.65	3153.71	26.89	231.63	6.98
LoginPage		1000	0	0.00%	1549.92	25	5182	1211.50	2786.80	4987.90	5104.85	23.72	227.43	3.57
Logout		1000	0	0.00%	1081.18	19	3938	976.00	1944.80	2221.80	3020.82	27.35	262.27	4.33




C. 500 Users




Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total	6000	6000	100.00%	77.94	0	1874	0.00	46.00	442.00	1587.98	720.29	1905.53	0.00
HomePage	1500	1500	100.00%	280.05	0	1874	1.00	1439.00	1574.95	1855.99	181.38	479.84	0.00
LoggedIn	1500	1500	100.00%	8.45	0	495	0.00	17.00	45.00	157.96	227.00	600.52	0.00
LoginPage	1500	1500	100.00%	16.35	0	627	0.00	42.00	110.95	442.99	227.03	600.61	0.00
Logout	1500	1500	100.00%	6.89	0	502	0.00	13.00	34.00	125.95	229.29	606.58	0.00




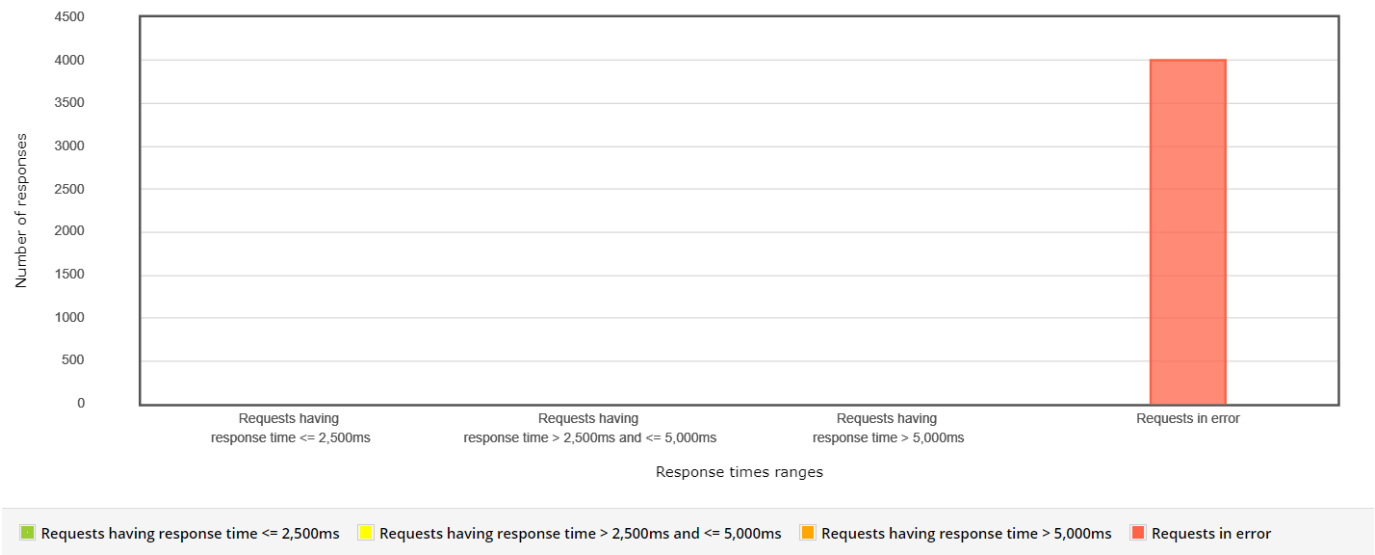
■ Requests having response time <= 2,500ms ■ Requests having response time > 2,500ms and <= 5,000ms ■ Requests having response time > 5,000ms ■ Requests in error



D. 1000 Users



Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label ^	#Samples ⇅	FAIL ⇅	Error % ⇅	Average ⇅	Min ⇅	Max ⇅	Median ⇅	90th pct ⇅	95th pct ⇅	99th pct ⇅	Transactions/s ⇅	Received ⇅	Sent ⇅
Total	4000	4000	100.00%	92810.41	0	373242	40.50	270602.80	287153.95	372619.99	10.66	36.65	0.33
HomePage	1000	1000	100.00%	142237.28	80038	373242	118722.50	189441.20	372541.60	373157.91	2.68	16.30	0.34
LogedIn	1000	1000	100.00%	6.91	0	267	0.00	11.90	41.90	179.62	430.29	1138.34	0.00
LoginPage	1000	1000	100.00%	228994.45	0	292992	254334.50	287326.90	291250.90	292933.98	3.39	8.07	0.00
Logout	1000	1000	100.00%	3.01	0	117	0.00	1.00	12.95	93.00	524.11	1386.53	0.00



VII. DISCUSSION

A. Time behavior

We can see that the current version of TeaStore is more performant than version 1.0.1. This information is obtained from the response timetables. Version 1.0.1 clearly runs faster in the 20 users benchmark, with average, maximum, and median being way lower than the same benchmark of the current version. However, for the 200 users benchmark, which represents the average load, the average response time of the current version is better, but not the maximum nor the median. The better performance of the current version is not apparent until we look at another presented graph, the one showing a total count of responses in 4 different response time categories. This graph clearly shows that a more significant ratio of requests takes less than 2.5 seconds compared to the older version. However, both versions could not pass the set objective of 99% response times of requests under 2.5 seconds, even for an average load. If we look at 500 and 1000 users, we can see TeaStore is clearly not able to keep up. Not only the average response times is exceptionally high, but these requests also fail. In conclusion, both versions did not pass the Time behavior metric.

B. Resource utilization

There is no apparent difference between the two versions. Both versions also passed the set objective. The only noteworthy characteristic we can read from the graphs is the slowly increasing usage of memory. However, this does not present a genuine issue either since TeaStore is written in Java which automatically garbage collects.

C. Capacity

Looking at the last graph of each benchmark, we can see both versions have serious issues keeping up with the above-average load. Additionally, the older version has issues even on the average load of 200 users. This clearly shows that both versions did not pass the Capacity metric.

REFERENCES

- [1] <https://github.com/DescartesResearch/TeaStore>
- [2] <https://www.ej-technologies.com/products/jprofiler/overview.html>